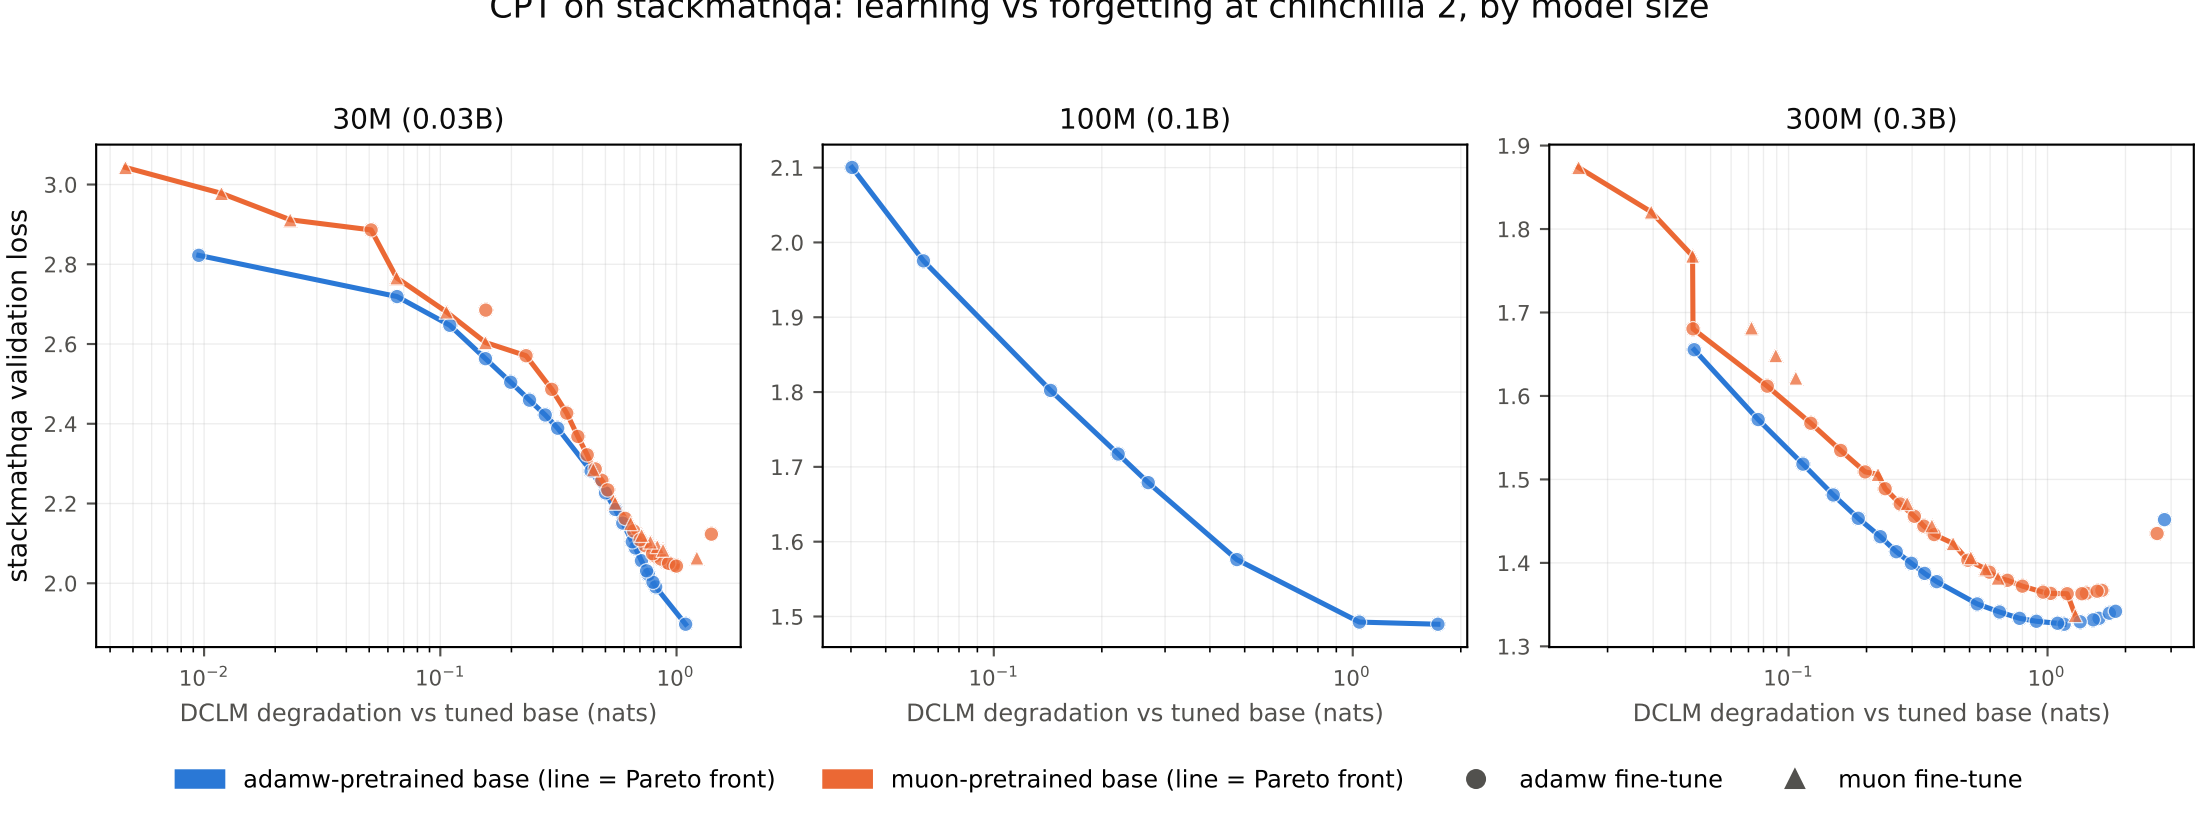
CPT on stackmathqa: learning vs forgetting at chinchilla 2, by model size
30M (0.03B)
100M (0.1B)
300M (0.3B)
stackmathqa validation loss
DCLM degradation vs tuned base (nats)
DCLM degradation vs tuned base (nats)
DCLM degradation vs tuned base (nats)
adamw-pretrained base (line = Pareto front)
muon-pretrained base (line = Pareto front)
adamw fine-tune
muon fine-tune