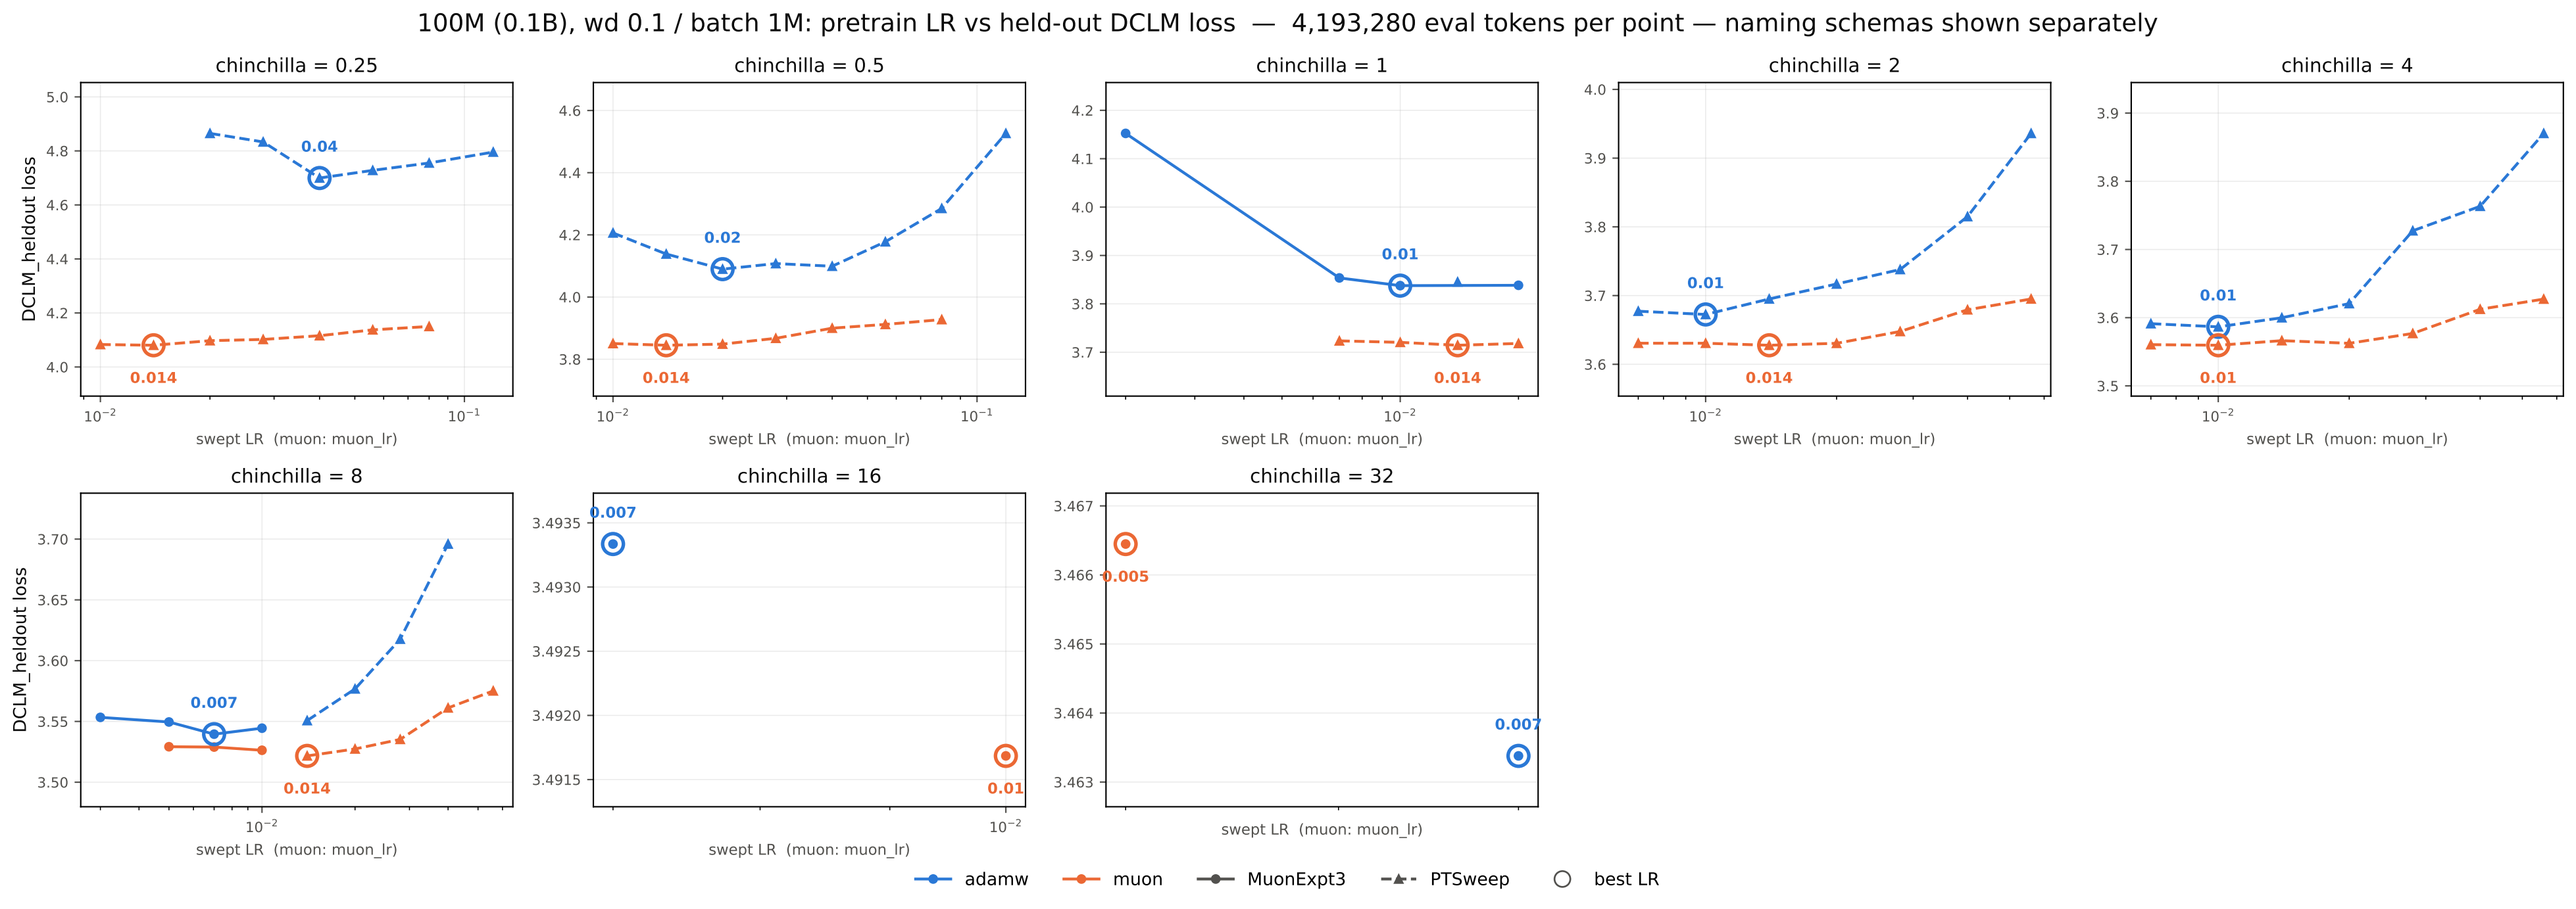
100M (0.1B), wd 0.1 / batch 1M: pretrain LR vs held-out DCLM loss — 4,193,280 eval tokens per point — naming schemas shown separately
chinchilla = 0.25
chinchilla = 0.5
chinchilla = 1
chinchilla = 2
chinchilla = 4
chinchilla = 8
chinchilla = 16
chinchilla = 32
DCLM_heldout loss
swept LR (muon: muon_lr)
0.04
0.014
0.02
0.01
0.007
0.005
adamw
muon
MuonExpt3
PTSweep
best LR