

100M (0.1B), wd 0.1 / batch 1M: pretrain LR vs held-out DCLM loss — 4,193,280 eval tokens per point — naming schemas shown separately

