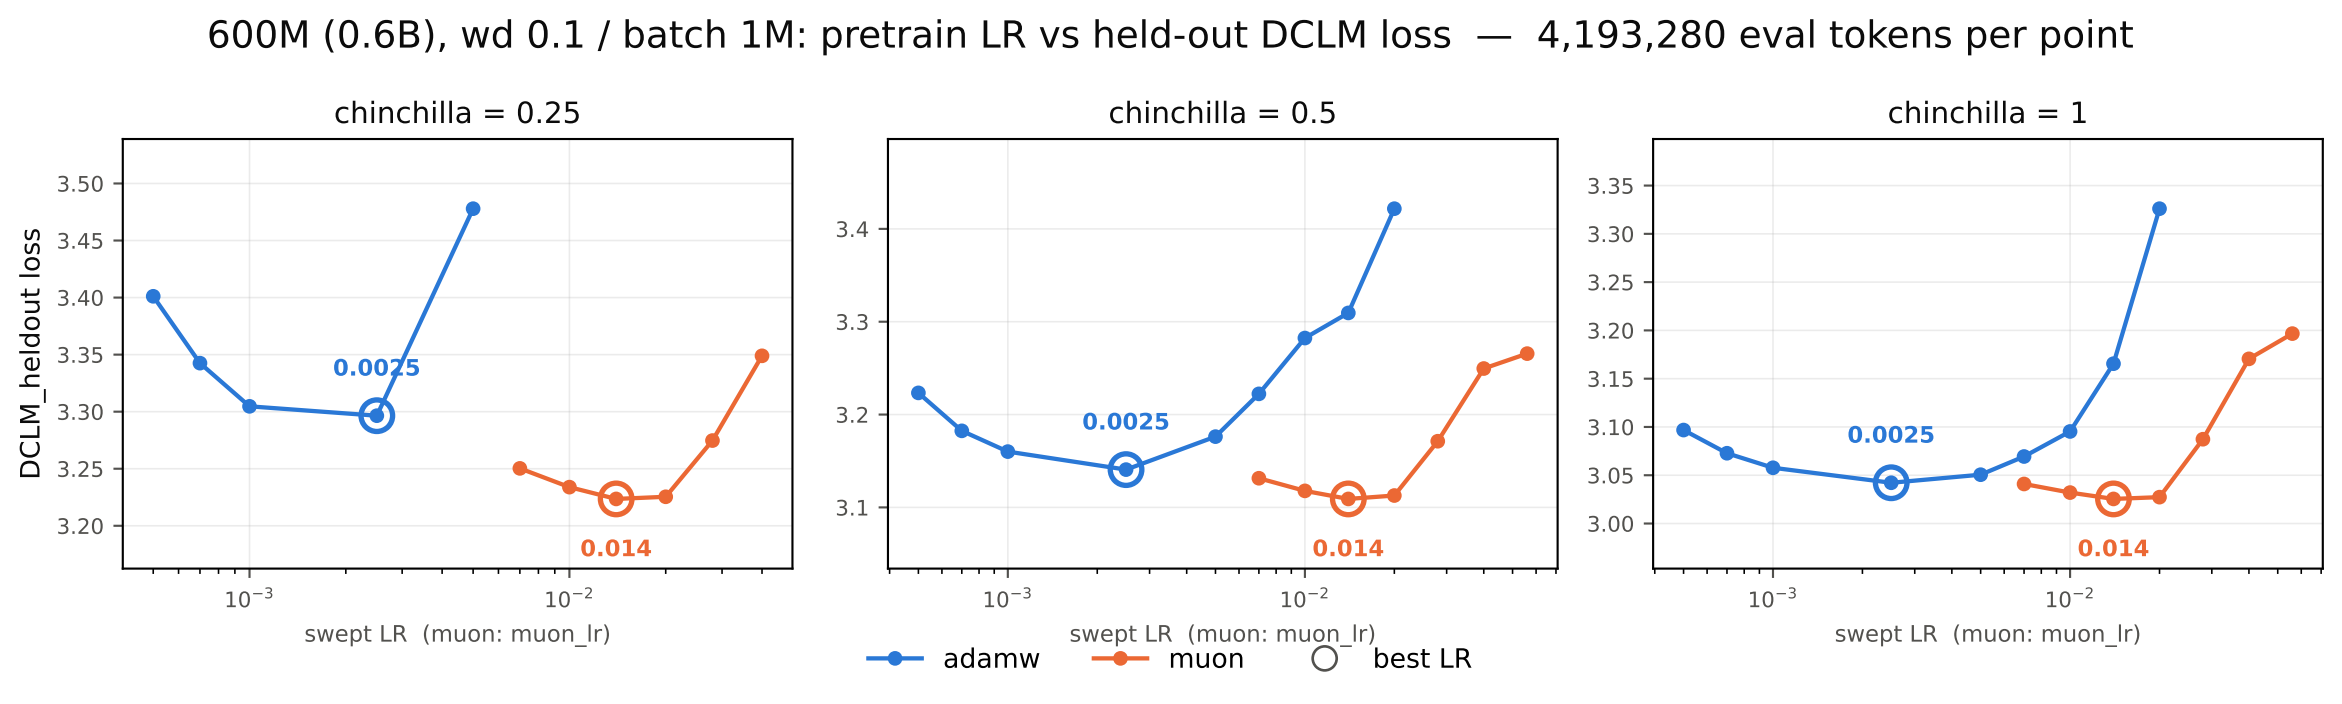
600M (0.6B), wd 0.1 / batch 1M: pretrain LR vs held-out DCLM loss — 4,193,280 eval tokens per point
chinchilla = 0.25
chinchilla = 0.5
chinchilla = 1
DCLM_heldout loss
swept LR (muon: muon_lr)
0.0025
0.014
adamw
muon
best LR