

600M (0.6B), wd 0.1 / batch 1M: pretrain LR vs held-out DCLM loss — 4,193,280 eval tokens per point

