

Best-LR held-out DCLM loss vs token budget (wd 0.1, batch 1M)

