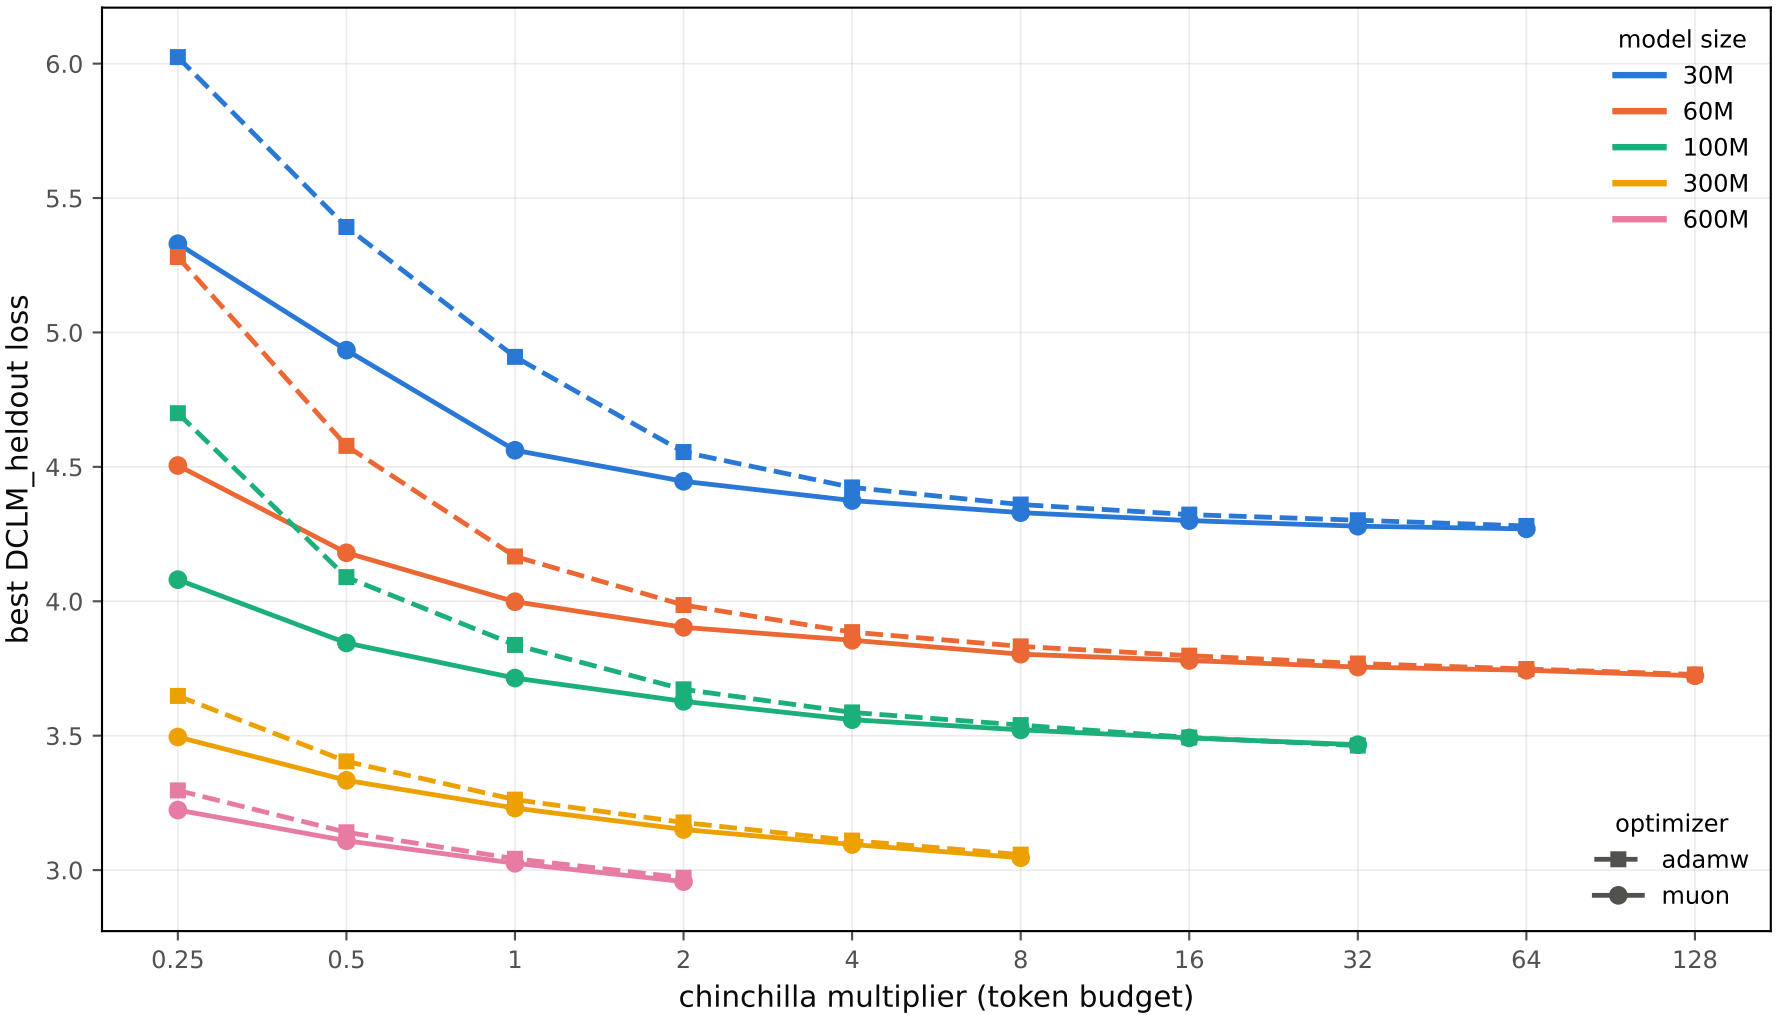
Best-LR held-out DCLM loss vs token budget, all model sizes
(wd 0.1, batch 1M)
best DCLM_heldout loss
6.0
5.5
5.0
4.5
4.0
3.5
3.0
0.25
0.5
1
2
4
8
16
32
64
128
chinchilla multiplier (token budget)
model size
30M
60M
100M
300M
600M
optimizer
adamw
muon