

Best-LR held-out DCLM loss vs token budget, all model sizes
(wd 0.1, batch 1M; annotation = the winning LR)

