

60M adamw: loss of the best-LR model vs token budget
(annotation = the LR that won that budget)

