

60M (0.06B), wd 0.1: pretrain LR vs held-out DCLM loss across global batch size — 4,193,280 eval tokens/point

