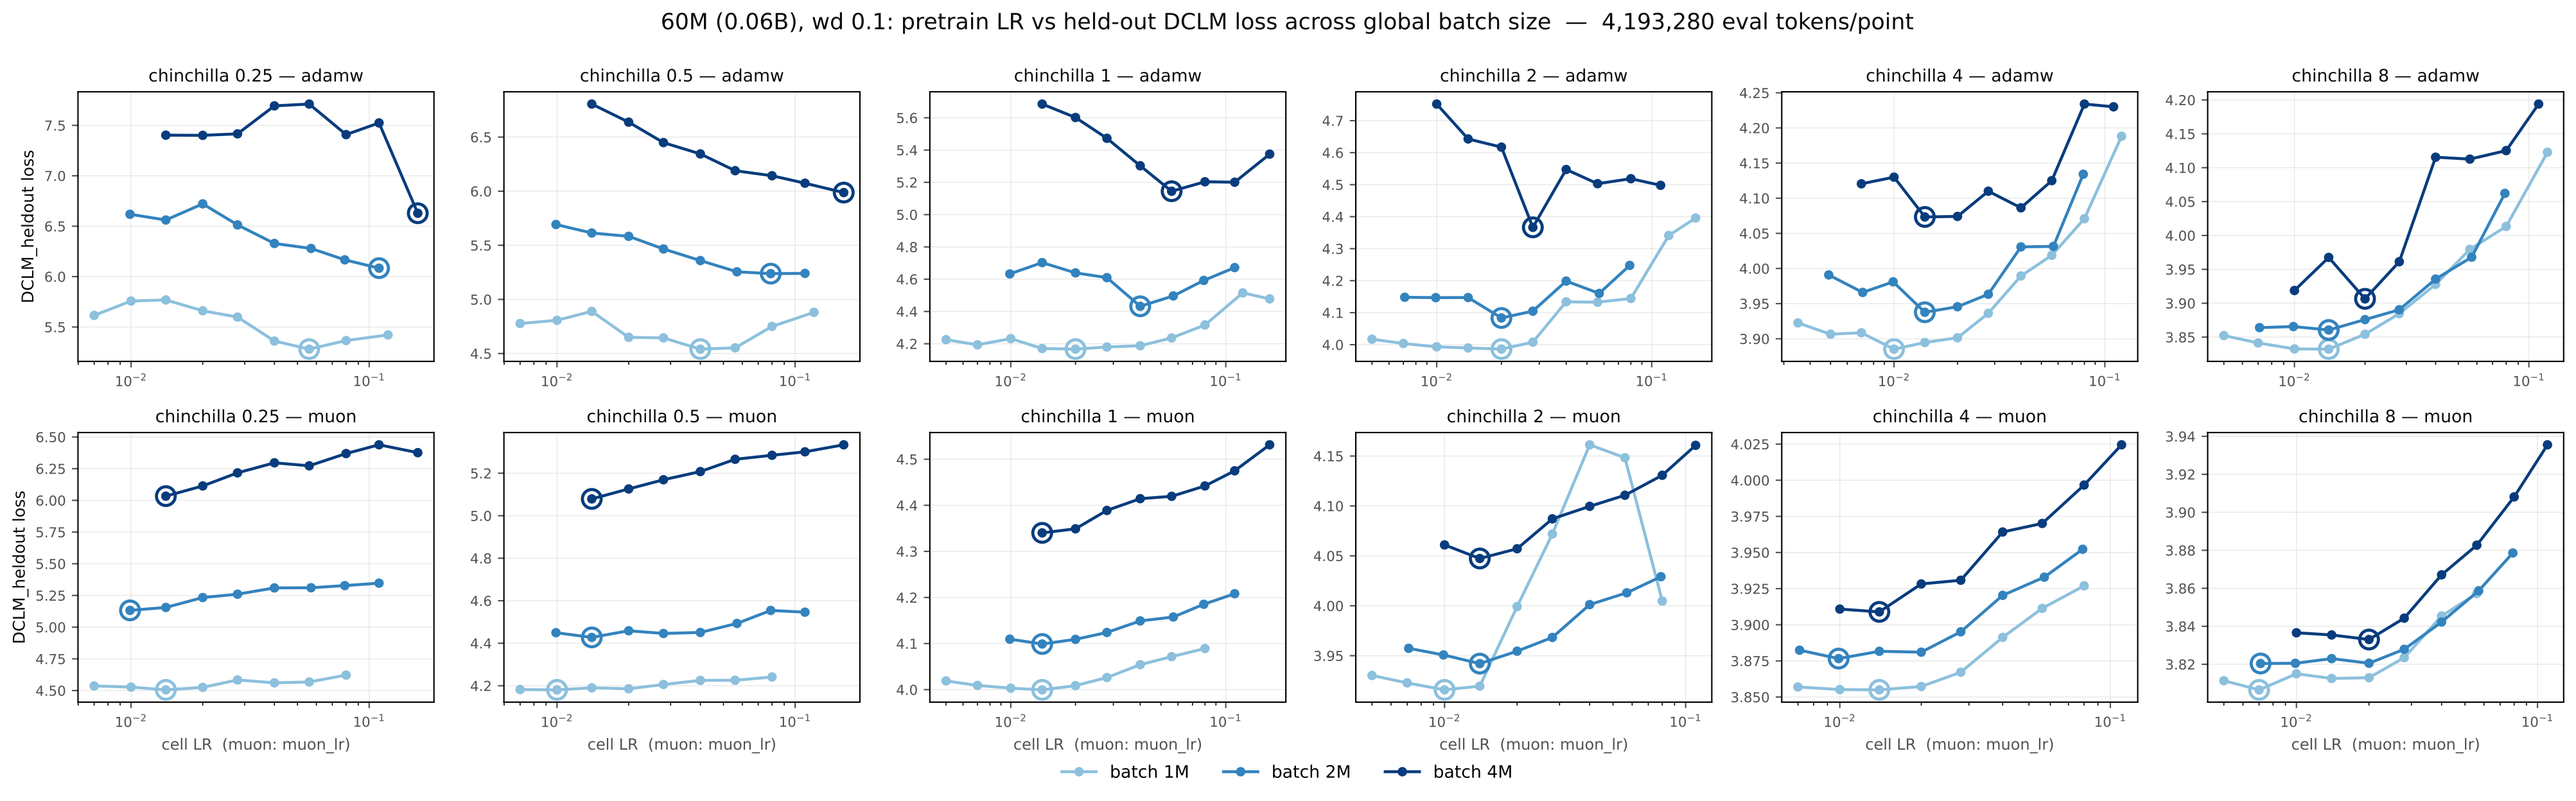
60M (0.06B), wd 0.1: pretrain LR vs held-out DCLM loss across global batch size — 4,193,280 eval tokens/point
chinchilla 0.25 — adamw
chinchilla 0.5 — adamw
chinchilla 1 — adamw
chinchilla 2 — adamw
chinchilla 4 — adamw
chinchilla 8 — adamw
chinchilla 0.25 — muon
chinchilla 0.5 — muon
chinchilla 1 — muon
chinchilla 2 — muon
chinchilla 4 — muon
chinchilla 8 — muon
DCLM_heldout loss
cell LR (muon: muon_lr)
batch 1M
batch 2M
batch 4M