

100M (0.1B) best pretrain LR by batch size and token budget
cell = best LR / its DCLM_heldout loss; row = chinchilla; wd 0.1

	adamw bs 1M	adamw bs 2M	adamw bs 4M	muon bs 1M	muon bs 2M	muon bs 4M
0.25	0.04 4.700	--	--	0.014 4.080	--	--
0.5	0.02 4.090	--	--	0.014 3.845	--	--
2	0.01 3.673	0.014 3.798	--	0.014 3.628	--	--
4	0.01 3.587	--	--	0.01 3.560	--	--
8	0.014 3.551	--	--	0.014 3.522	--	--