

600M (0.6B) best pretrain LR by batch size and token budget
cell = best LR / its DCLM_heldout loss; row = chinchilla; wd 0.1

	adamw bs 1M	adamw bs 2M	adamw bs 4M	muon bs 1M	muon bs 2M	muon bs 4M
0.25	0.0025 3.296	--	--	0.014 3.224	--	--
0.5	0.0025 3.141	--	--	0.014 3.109	--	--
1	0.0025 3.042	--	--	0.014 3.025	--	--