

600M (0.6B) best pretrain LR by batch size and token budget
cell = best LR / its DCLM_heldout loss; row = chinchilla; wd 0.1

	adamw bs 1M	adamw bs 2M	adamw bs 4M	muon bs 1M	muon bs 2M	muon bs 4M
0.5	0.005 3.176	--	--	0.056 3.266	--	--
1	0.005 3.051	--	--	0.056 3.197	--	--