

Pretrain LR vs held-out DCLM loss across global batch size, every size x token budget (wd 0.1) — 4,193,280 eval tokens/point

