

60M adamw (wd 0.1, batch 1M): pretrain LR vs held-out DCLM loss — 4,193,280 eval tokens per point

