


Pretrain LR vs held-out DCLM loss, every size x token budget (wd 0.1, batch 1M) — 4,193,280 eval tokens per point



adamw muon ○ table LR (annotated) ★ newly trained (eval written 2026-09-17 16:00Z-19:18Z)