


Pretrain LR vs held-out DCLM loss, every size x token budget (wd 0.1, batch 1M) — 4,193,280 eval tokens per point



adamw muon ○ table LR (annotated)