


Pretrain LR vs held-out DCLM loss, every size x token budget (wd 0.1, batch 1M) — 4,193,280 eval tokens per point



—●— adamw    —●— muon    ○ table LR (annotated)    ★ newly trained (since 2026-09-17)